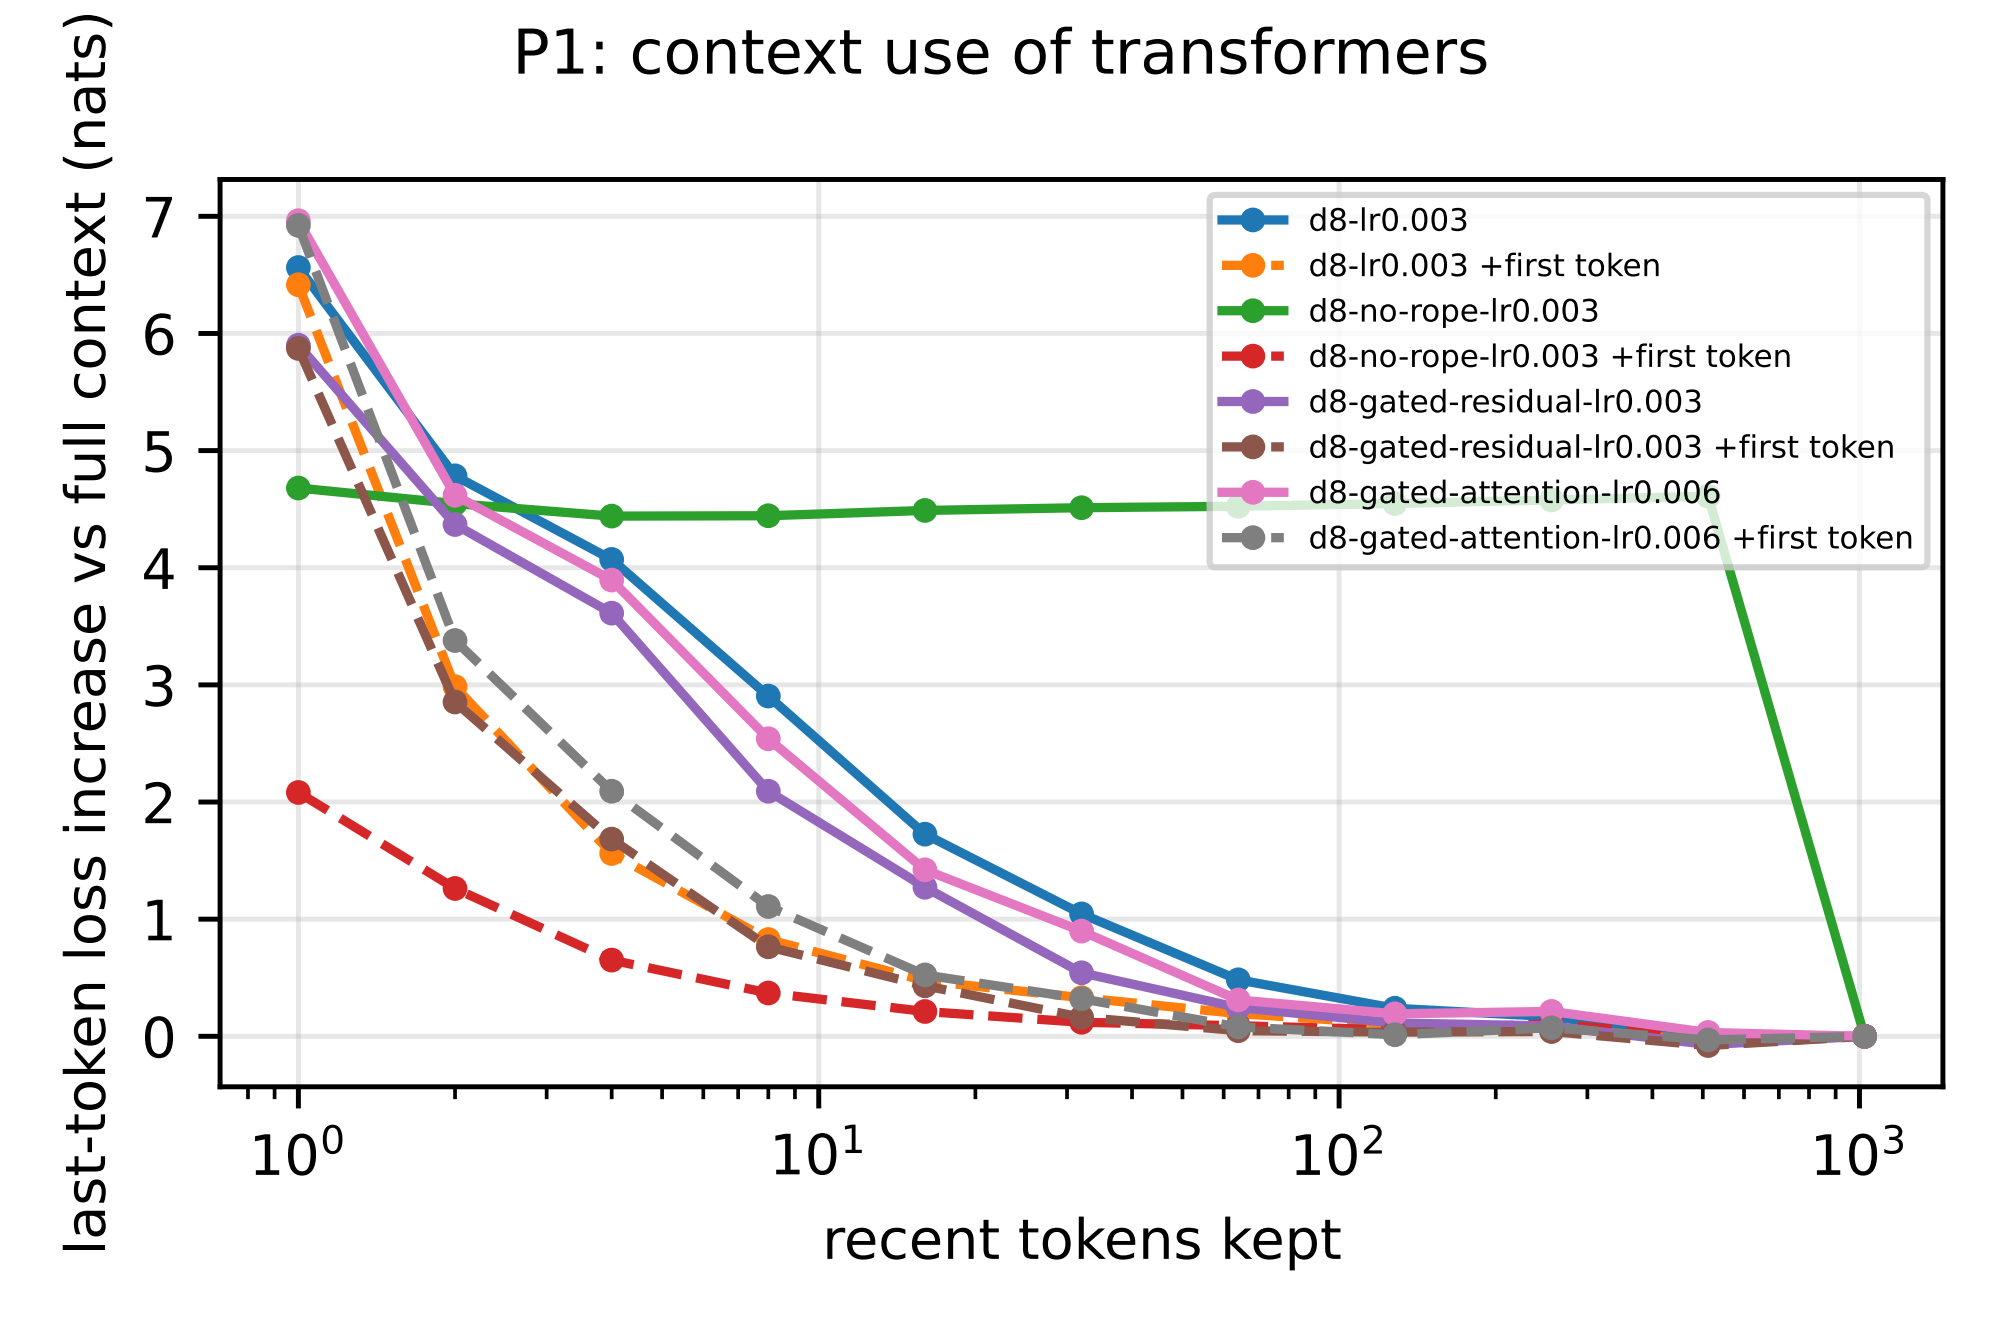
P1: context use of transformers
last-token loss increase vs full context (nats)
recent tokens kept
d8-lr0.003
d8-lr0.003 +first token
d8-no-rope-lr0.003
d8-no-rope-lr0.003 +first token
d8-gated-residual-lr0.003
d8-gated-residual-lr0.003 +first token
d8-gated-attention-lr0.006
d8-gated-attention-lr0.006 +first token