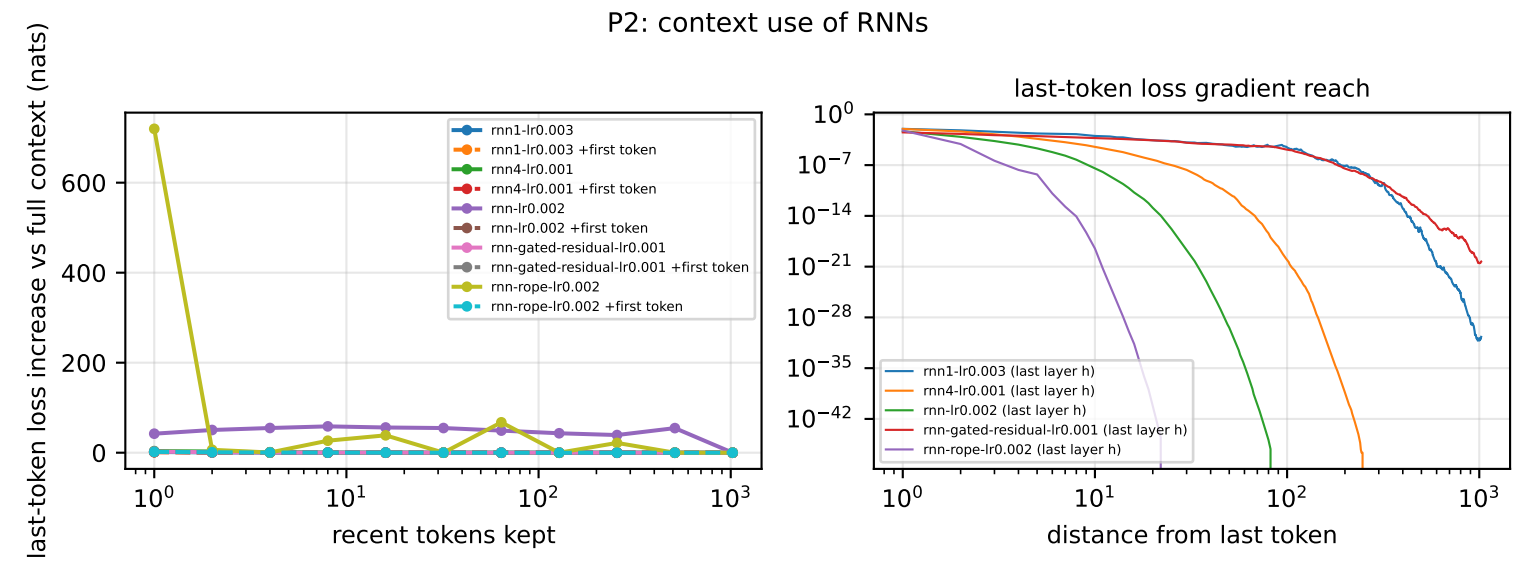
P2: context use of RNNs
last-token loss increase vs full context (nats)
recent tokens kept
rnn1-lr0.003
rnn1-lr0.003 +first token
rnn4-lr0.001
rnn4-lr0.001 +first token
rnn-lr0.002
rnn-lr0.002 +first token
rnn-gated-residual-lr0.001
rnn-gated-residual-lr0.001 +first token
rnn-rope-lr0.002
rnn-rope-lr0.002 +first token
last-token loss gradient reach
distance from last token
rnn1-lr0.003 (last layer h)
rnn4-lr0.001 (last layer h)
rnn-lr0.002 (last layer h)
rnn-gated-residual-lr0.001 (last layer h)
rnn-rope-lr0.002 (last layer h)