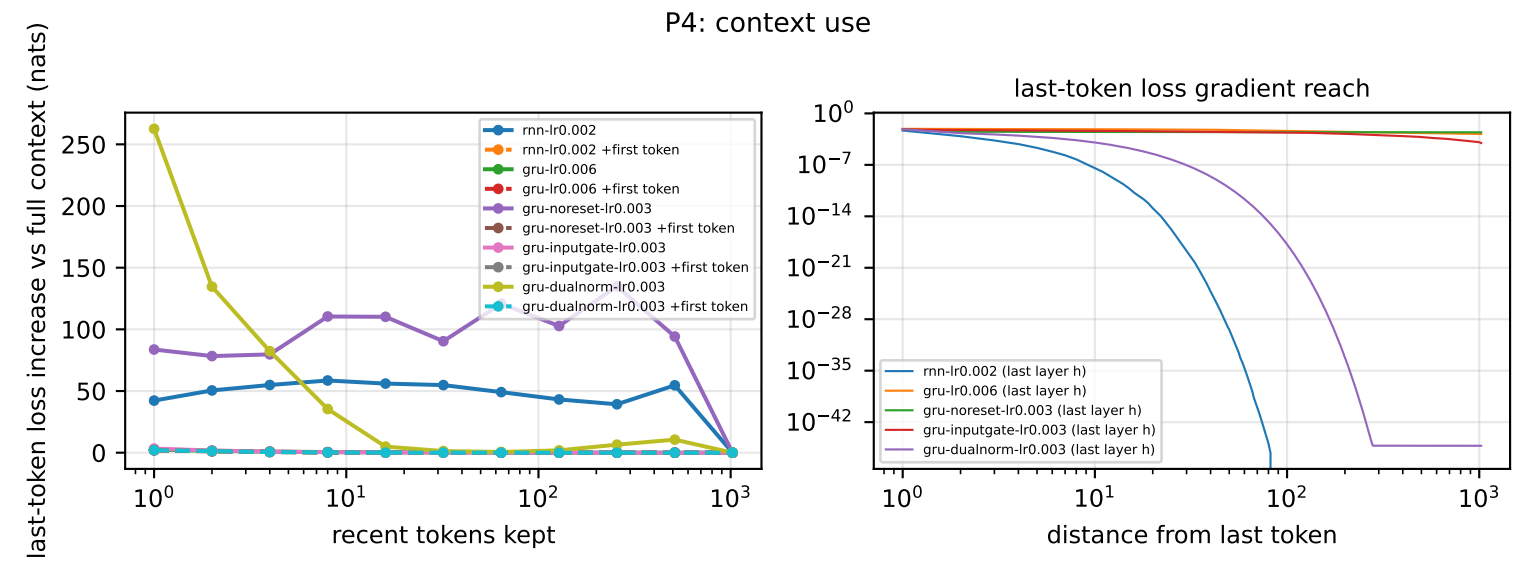
P4: context use
last-token loss increase vs full context (nats)
recent tokens kept
rnn-lr0.002
rnn-lr0.002 +first token
gru-lr0.006
gru-lr0.006 +first token
gru-noreset-lr0.003
gru-noreset-lr0.003 +first token
gru-inputgate-lr0.003
gru-inputgate-lr0.003 +first token
gru-dualnorm-lr0.003
gru-dualnorm-lr0.003 +first token
last-token loss gradient reach
distance from last token
rnn-lr0.002 (last layer h)
gru-lr0.006 (last layer h)
gru-noreset-lr0.003 (last layer h)
gru-inputgate-lr0.003 (last layer h)
gru-dualnorm-lr0.003 (last layer h)